

| |
|---|
| Experiment No. 3 |
| Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model |
| Date of Performance: |
| Date of Submission: |

Aim:

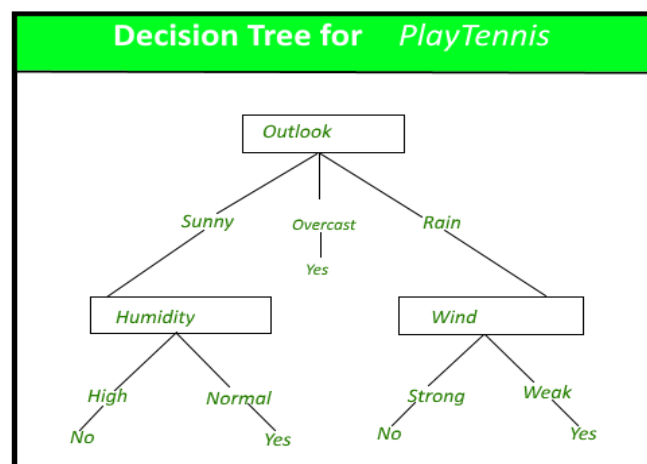
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective:

To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

```
import os
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
%matplotlib inline
```

```
# To ignore warning messages
```

```
import warnings
```

```
warnings.filterwarnings('ignore')
```

```
# Adult dataset path
```

```
adult_dataset_path = "../input/adult_dataset.csv"
```

```
# Function for loading adult dataset
```

```
def load_adult_data(adult_path=adult_dataset_path):
```

```
    csv_path = os.path.join(adult_path)
```

```
    return pd.read_csv(csv_path)
```

```
# Calling load adult function and assigning to a new variable df
```

```
df = load_adult_data()
```

```
# load top 3 rows values from adult dataset
```

```
df.head(3)
```

```
:
```

| | age | workclass | fnlwgt | education | education.num | marital.status | occupation | relationship | race | sex | capital.gain |
|---|-----|-----------|--------|--------------|---------------|----------------|-----------------|---------------|-------|--------|--------------|
| 0 | 90 | ? | 77053 | HS-grad | 9 | Widowed | ? | Not-in-family | White | Female | 0 |
| 1 | 82 | Private | 132870 | HS-grad | 9 | Widowed | Exec-managerial | Not-in-family | White | Female | 0 |
| 2 | 66 | ? | 186061 | Some-college | 10 | Widowed | ? | Unmarried | Black | Female | 0 |

```
print ("Rows : ",df.shape[0])
```

```
print ("Columns : ",df.shape[1])
```

```
print ("\nFeatures : \n",df.columns.tolist())
```

```
print ("\nMissing values : ", df.isnull().sum().values.sum())
```

```
print ("\nUnique values : \n",df.nunique())
```

`df.describe()`

| | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week |
|-------|--------------|--------------|---------------|--------------|--------------|----------------|
| count | 32561.000000 | 3.256100e+04 | 32561.000000 | 32561.000000 | 32561.000000 | 32561.000000 |
| mean | 38.581647 | 1.897784e+05 | 10.080679 | 1077.648844 | 87.303830 | 40.437456 |
| std | 13.640433 | 1.055500e+05 | 2.572720 | 7385.292085 | 402.960219 | 12.347429 |
| min | 17.000000 | 1.228500e+04 | 1.000000 | 0.000000 | 0.000000 | 1.000000 |
| 25% | 28.000000 | 1.178270e+05 | 9.000000 | 0.000000 | 0.000000 | 40.000000 |
| 50% | 37.000000 | 1.783560e+05 | 10.000000 | 0.000000 | 0.000000 | 40.000000 |
| 75% | 48.000000 | 2.370510e+05 | 12.000000 | 0.000000 | 0.000000 | 45.000000 |
| max | 90.000000 | 1.484705e+06 | 16.000000 | 99999.000000 | 4356.000000 | 99.000000 |

pull top 5 row values to understand the data and how it's look like

`df.head()`

checking "?" total values present in particular 'workclass' feature

`df_check_missing_workclass = (df['workclass']=='?').sum()`

`df_check_missing_workclass`

.....

| | age | workclass | fnlwgt | education | education.num | marital.status | occupation | relationship | race | sex | capital.gain |
|---|-----|-----------|--------|--------------|---------------|----------------|-------------------|---------------|-------|--------|--------------|
| 0 | 90 | ? | 77053 | HS-grad | 9 | Widowed | ? | Not-in-family | White | Female | 0 |
| 1 | 82 | Private | 132870 | HS-grad | 9 | Widowed | Exec-managerial | Not-in-family | White | Female | 0 |
| 2 | 66 | ? | 186061 | Some-college | 10 | Widowed | ? | Unmarried | Black | Female | 0 |
| 3 | 54 | Private | 140359 | 7th-8th | 4 | Divorced | Machine-op-inspct | Unmarried | White | Female | 0 |
| 4 | 41 | Private | 264663 | Some-college | 10 | Separated | Prof-specialty | Own-child | White | Female | 0 |

checking "?" total values present in particular 'occupation' feature

`df_check_missing_occupation = (df['occupation']=='?').sum()`

`df_check_missing_occupation`

```

age          0
workclass    1836
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   1843
relationship 0
race         0
sex          0
capital.gain 0
capital.loss 0
hours.per.week 0
native.country 583
income       0
dtype: int64

```

```
percent_missing = (df=="?").sum() * 100/len(df)
```

```
percent_missing
```

```

age          0.000000
workclass     5.638647
fnlwgt        0.000000
education     0.000000
education.num 0.000000
marital.status 0.000000
occupation    5.660146
relationship   0.000000
race          0.000000
sex           0.000000
capital.gain   0.000000
capital.loss   0.000000
hours.per.week 0.000000
native.country 1.790486
income        0.000000
dtype: float64

```

contain any missing value as '?'

```
df.apply(lambda x: x != '?',axis=1).sum()
```

```
age          32561
workclass     30725
fnlwgt        32561
education     32561
education.num 32561
marital.status 32561
occupation    30718
relationship   32561
race          32561
sex           32561
capital.gain   32561
capital.loss   32561
hours.per.week 32561
native.country 31978
income        32561
dtype: int64
```

```
df = df[df['workclass'] != '?']
```

```
df.head()
```

| | age | workclass | fnlwgt | education | education.num | marital.status | occupation | relationship | race | sex | capital.gain |
|---|-----|-----------|--------|--------------|---------------|----------------|-------------------|---------------|-------|--------|--------------|
| 1 | 82 | Private | 132870 | HS-grad | 9 | Widowed | Exec-managerial | Not-in-family | White | Female | 0 |
| 3 | 54 | Private | 140359 | 7th-8th | 4 | Divorced | Machine-op-inspct | Unmarried | White | Female | 0 |
| 4 | 41 | Private | 264663 | Some-college | 10 | Separated | Prof-specialty | Own-child | White | Female | 0 |
| 5 | 34 | Private | 216864 | HS-grad | 9 | Divorced | Other-service | Unmarried | White | Female | 0 |
| 6 | 38 | Private | 150601 | 10th | 6 | Separated | Adm-clerical | Unmarried | White | Male | 0 |

```
df_categorical = df.select_dtypes(include=['object'])
```

```
df_categorical.apply(lambda x: x == '?',axis=1).sum()
```

```
df = df[df['occupation'] != '?']
```



```
df = df[df['native.country'] != '?']
```

```
# dropping the "?"s from occupation and native.country
```

```
df = df[df['occupation'] != '?']
```

```
df = df[df['native.country'] != '?']
```

```
workclass      0
education      0
marital.status  0
occupation     7
relationship   0
race           0
sex           0
native.country 556
income         0
dtype: int64
```

```
# check the dataset whether cleaned or not?
```

```
df.info()
```

```
from sklearn import preprocessing
```

```
# encode categorical variables using label Encoder
```

```
# select all categorical variables
```

```
df_categorical = df.select_dtypes(include=['object'])
```

```
df_categorical.head()
```

| | workclass | education | marital.status | occupation | relationship | race | sex | native.country | income |
|---|-----------|--------------|----------------|-------------------|---------------|-------|--------|----------------|--------|
| 1 | Private | HS-grad | Widowed | Exec-managerial | Not-in-family | White | Female | United-States | <=50K |
| 3 | Private | 7th-8th | Divorced | Machine-op-inspct | Unmarried | White | Female | United-States | <=50K |
| 4 | Private | Some-college | Separated | Prof-specialty | Own-child | White | Female | United-States | <=50K |
| 5 | Private | HS-grad | Divorced | Other-service | Unmarried | White | Female | United-States | <=50K |
| 6 | Private | 10th | Separated | Adm-clerical | Unmarried | White | Male | United-States | <=50K |

```
le = preprocessing.LabelEncoder()
```

```
df_categorical = df_categorical.apply(le.fit_transform)
```

```
df_categorical.head()
```

| | workclass | education | marital.status | occupation | relationship | race | sex | native.country | income |
|---|-----------|-----------|----------------|------------|--------------|------|-----|----------------|--------|
| 1 | 2 | 11 | 6 | 3 | 1 | 4 | 0 | 38 | 0 |
| 3 | 2 | 5 | 0 | 6 | 4 | 4 | 0 | 38 | 0 |
| 4 | 2 | 15 | 5 | 9 | 3 | 4 | 0 | 38 | 0 |
| 5 | 2 | 11 | 0 | 7 | 4 | 4 | 0 | 38 | 0 |
| 6 | 2 | 0 | 5 | 0 | 4 | 4 | 1 | 38 | 0 |

```
df = df.drop(df_categorical.columns,axis=1)
```

```
df = pd.concat([df,df_categorical],axis=1)
```

```
df.head()
```

```
df.info()
```

```
df['income'] = df['income'].astype('category')
```

| | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | workclass | education | marital.status | occupation |
|---|-----|--------|---------------|--------------|--------------|----------------|-----------|-----------|----------------|------------|
| 1 | 82 | 132870 | 9 | 0 | 4356 | 18 | 2 | 11 | 6 | 3 |
| 3 | 54 | 140359 | 4 | 0 | 3900 | 40 | 2 | 5 | 0 | 6 |
| 4 | 41 | 264663 | 10 | 0 | 3900 | 40 | 2 | 15 | 5 | 9 |
| 5 | 34 | 216864 | 9 | 0 | 3770 | 45 | 2 | 11 | 0 | 7 |
| 6 | 38 | 150601 | 6 | 0 | 3770 | 40 | 2 | 0 | 5 | 0 |

```
df.info()
```

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop('income',axis=1)
```

```
# Putting response/dependent variable/feature to y
```

```
y = df['income']
```

```
X.head(3)
```

| | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | workclass | education | marital.status | occupation |
|---|-----|--------|---------------|--------------|--------------|----------------|-----------|-----------|----------------|------------|
| 1 | 82 | 132870 | 9 | 0 | 4356 | 18 | 2 | 11 | 6 | 3 |
| 3 | 54 | 140359 | 4 | 0 | 3900 | 40 | 2 | 5 | 0 | 6 |
| 4 | 41 | 264663 | 10 | 0 | 3900 | 40 | 2 | 15 | 5 | 9 |

```
y.head(3)
```

```
1    0
3    0
4    0
Name: income, dtype: category
Categories (2, int64): [0, 1]
```

```
# Splitting the data into train and test
```

```
X_train,X_test,y_train,y_test
```

=

```
train_test_split(X,y,test_size=0.30,random_state=99)
```

```
X_train.head()
```

| | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | workclass | education | marital.status | occupa |
|-------|-----|--------|---------------|--------------|--------------|----------------|-----------|-----------|----------------|--------|
| 24351 | 42 | 289636 | 9 | 0 | 0 | 46 | 2 | 11 | 2 | 13 |
| 15626 | 37 | 52465 | 9 | 0 | 0 | 40 | 1 | 11 | 4 | 7 |
| 4347 | 38 | 125933 | 14 | 0 | 0 | 40 | 0 | 12 | 2 | 9 |
| 23972 | 44 | 183829 | 13 | 0 | 0 | 38 | 5 | 9 | 4 | 0 |
| 26843 | 35 | 198841 | 11 | 0 | 0 | 35 | 2 | 8 | 0 | 12 |

```
from sklearn.tree import DecisionTreeClassifier
```

```
# Fitting the decision tree with default hyperparameters, apart from
```

```
# max_depth which is 5 so that we can plot and read the tree.
```

```
dt_default = DecisionTreeClassifier(max_depth=5)
```

```
dt_default.fit(X_train,y_train)
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,  
                        max_features=None, max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, presort=False,  
                        random_state=None, splitter='best')
```

```
From sklearn metricimport
```

```
classification_report,confusion_matrix,accuracy_score
```

```
# making predictions
```

```
y_pred_default = dt_default.predict(X_test)
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.86 | 0.95 | 0.91 | 6867 |
| 1 | 0.78 | 0.52 | 0.63 | 2182 |
| accuracy | | | 0.85 | 9049 |
| macro avg | 0.82 | 0.74 | 0.77 | 9049 |
| weighted avg | 0.84 | 0.85 | 0.84 | 9049 |

```
print(confusion_matrix(y_test,y_pred_default))
```

```
print(accuracy_score(y_test,y_pred_default))
```

```
[[ 6553   314]
 [ 1038  1144]]
0.8505912255497845
```

```
!pip install pydotplus
```

```
from IPython.display import Image
```

```
from sklearn.externals.six import StringIO
```

```
from sklearn.tree import export_graphviz
```

```
import pydotplus,graphviz
```

```
# Putting features
```

```
features = list(df.columns[1:])
```

```
features
```

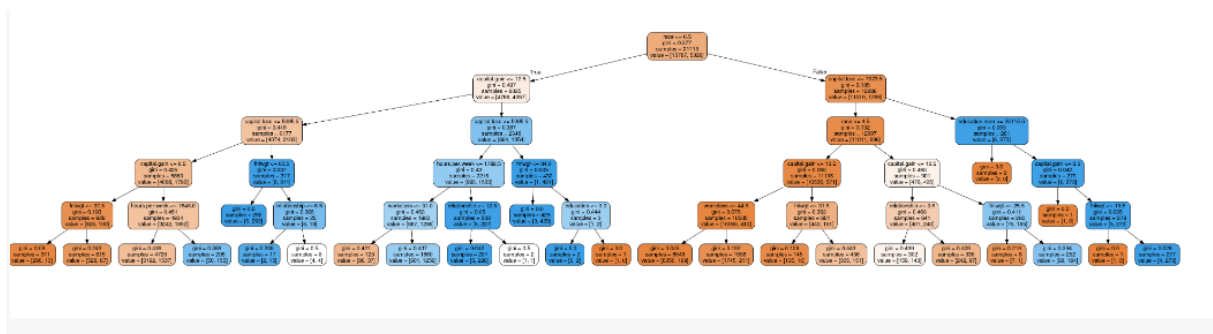
```
dot_data = StringIO()
```

```
export_graphviz(dt_default, out_file=dot_data,
```

```
feature_names=features, filled=True,rounded=True)
```

```
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
```

```
Image(graph.create_png())
```



Conclusion:

1. Discuss about the how categorical attributes have been dealt with during data pre-processing.

➤ Categorical attribute handling methods include label encoding, one-hot encoding, frequency encoding, target encoding, binary encoding, embedding layers (for neural networks), and handling missing values appropriately.

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.

- Hyperparameter tuning involves adjusting parameters like `max_depth`, `min_samples_leaf`, `min_samples_split`, `criterion`, `max_features`, `min_impurity_decrease`, `ccp_alpha`, and others.
- The tuning process aims to find the optimal combination of hyperparameters that balances model complexity and predictive performance.
- Techniques like Grid Search or Random Search can be used for systematic hyperparameter exploration.
- Careful tuning can improve the Decision Tree's accuracy and prevent overfitting or underfitting.

3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.96 | 0.90 | 6867 |
| 1 | 0.77 | 0.47 | 0.59 | 2182 |
| accuracy | | | 0.84 | 9049 |
| macro avg | 0.81 | 0.71 | 0.74 | 9049 |
| weighted avg | 0.83 | 0.84 | 0.82 | 9049 |

